

COSC 364

RIP Assignment Report

29 April 2025

Jordan Chubb 62646119
Wenlong Zong 42718910

Percentage contribution

- Jordan Chubb - 60%
- Wenlong Zong - 40%

List of contributions

Jordan Chubb

- Argument parsing
- Configuration file processing
- Socket setup
- RipManager and RoutingTableEntry classes
- Automatic testing

Wenlong Zong

- Timers
- Packet building
- Manual testing
- Dijkstra's algorithm (automatic testing)

Test discussion

Configuration Files

Correctness of configuration files is checked thoroughly in a specific module. Functions within this module are tested using the doctest module. This module can also test the validity of multiple configuration files at once, ensuring that they are all compatible with each other. The test checks that router-id's are unique and are between 1 and 64000, all port numbers are between 1024 and 64000, and metrics are between 1 and 16. Each output of every configuration file is checked that it matches up to exactly one other configuration file, and that the metrics between them are the same. For automatic testing, all the automatically generated configuration files are checked to be valid. The expected outcome of invalid configuration files is that the program will fail to launch and that an appropriate error message will be displayed, which is what was found during testing.

Manual Testing

Configuration files were generated for the example network in the assignment specification. It was expected that the routing tables would converge to the exact topology in the provided example. The outcome was that they converged correctly.

Different routers were then stopped with the expectation that routers who had the stopped-router as a next-hop would then timeout this route and set the metric to 16. This metric of 16 would then propagate to the other routes, and all routers who had the metric at 16 would start the deletion process before deleting the route from their table. The outcome was that this was exactly the case.

Automatic Testing

The automatic testing program works by launching multiple router daemons as individual processes which are then monitored. The network topology is randomly generated by the program, configuration files are made and saved based on this, and then the daemons are started. The

program monitors the daemon processes through their stdout, which uses a json-encoded list to communicate their entire routing table. For each router daemon, a minimum spanning tree is generated based on the previously generated network topology using Dijkstra's algorithm. The minimum spanning tree is compared against the current routers routing table to ensure that the routing table costs are equal to the minimum costs from Dijkstra's algorithm.

The automatic tests include:

- 100 routers, fully-connected (each router is connected to every other router)
- 100 routers, sparsely-connected (each router connects to one other router)

For each of these tests, the program waits until the network has converged correctly, and then once it has, it changes the topology by stopping (crashing) or starting 50 random routers, and then waits for convergence again. It repeats this stopping/starting step several times, waiting for convergence after each change. For these tests we expect every router to eventually converge. Our test results found that the network would converge every time for all tests. This showed that our daemon was correctly converging after topology changes, even for very large networks with groups of routers being disconnected from other groups of routers.

During the convergence process of automatic testing it was discovered that convergence would take over 1 minute after stopping/starting random routers and that the CPU usage was 100% for a 30 second period. It was considered that this may have been caused by the routers waiting for their routes to timeout before updating them. In section 3.9.2 of RIP RFC specification, there is an optional heuristic that switches to a new route if the metric is the same and it is at least halfway to expiry. This is a potential method to improve the long convergence time. This optional heuristic was implemented, but it did not improve the problem.

The problem was then considered to be caused by a large number of triggered updates being sent by routers, which in turn caused other routers to also send triggered updates, making the problem worse. Section 3.10.1 of RIP RFC specification was then implemented, and the outcome was that the CPU usage would no longer go to 100% during the convergence process. This showed that the triggered updates were the source of this problem.

Packet Testing

Malformed packets are randomly generated to check that all packets fit the message format. The RIP packets are expected with command number 2, version number 2, must be zero section for 2 bytes and followed with payload of 1 to 25 RIP entries. Each RIP entry should have address family identifier AF_INET(2), must be zero section for 2 bytes, destination IPv4 address, 2 must be zero sections and the current metric for the destination. All malformed packets are expected to be dropped and that is achieved by the program.

Example configuration file

File: /csse/users/jch442/COSC364/configuration_files/router1.ini

Page 1 of 1

```
[SETTINGS]
router-id = 1
input-ports = 1024, 1025, 1026
outputs = 2001-1-2, 7001-8-7, 6001-5-6
```

Source code

```
1  '''COSC364 RIP Assignment
2  Jordan Chubb
3  Vincent Zong
4  29/04/2025
5  '''
6
7  import argparse
8  from functools import partial
9  import json
10 import selectors
11 import socket
12 import time
13
14 from configmanager import read_config_file
15 from ripmanager import RipManager
16
17
18 MAX_PACKET_SIZE = 4 + 20 * 25 # header + rip entry * max number of rip entries
19
20
21 parser = argparse.ArgumentParser()
22 parser.add_argument("config", help="filename of the configuration file")
23 parser.add_argument("-d", "--debug", help="print debugging information", action="store_true")
24 parser.add_argument("--autotesting", help="for automatic testing", action="store_true")
25 args = parser.parse_args()
26
27
28 if args.autotesting:
29     """Force all print calls to flush immediately. Required for
30     automatic testing's reading of stdout.
31     """
32     print = partial(print, flush=True)
33
34
35 def main():
36     config = read_config_file(args.config)
37     debug(config)
38
39     sockets = get_sockets(config)
40     selector = selectors.DefaultSelector()
41     for sock in sockets:
42         selector.register(sock, selectors.EVENT_READ)
43
44     rip = RipManager(debug, config, sockets[0])
45
46     next_print_time = time.time()
47     while True:
48         next_print = max(0, next_print_time - time.time())
49         next_timeout = min(next_print, rip.next_timeout())
```

```
50     events = selector.select(timeout=next_timeout)
51
52     for key, _ in events:
53         sock = key.fileobj
54         message = sock.recv(MAX_PACKET_SIZE)
55         rip.incoming_message(message)
56         rip.send_any_updates()
57
58     if time.time() >= next_print_time:
59         next_print_time = time.time() + 1
60         if args.autotesting:
61             print(json.dumps(rip.table_list()))
62         else:
63             print(rip)
64
65
66 def debug(line):
67     if args.debug:
68         print(line)
69
70
71 def get_sockets(config):
72     """Return a socket for each input port."""
73     sockets = []
74     for port in config.input_ports:
75         sock = socket.socket(socket.AF_INET, socket.SOCK_DGRAM)
76         sock.bind(('127.0.0.1', port))
77         sockets.append(sock)
78     return sockets
79
80
81 if __name__ == "__main__":
82     main()
83
```

```
1 import math
2 import random
3 import time
4
5 from configmanager import routerid_is_valid, metric_is_valid
6
7
8 TIME_MULTIPLIER = 6
9
10 PERIODIC_UPDATE_DELAY =      30 / TIME_MULTIPLIER
11 TRIGGERED_UPDATE_DELAY =     5 / TIME_MULTIPLIER
12 ENTRY_TIMEOUT_DELAY =       180 / TIME_MULTIPLIER
13 GARBAGE_COLLECTION_DELAY =  120 / TIME_MULTIPLIER
14
15
16 INFINITE_METRIC = 16
17
18 POISONED_REVERSE = True
19
20
21 class RipManager:
22     """This class manages the Routing Information Protocol. The routing
23         table is a dictionary where the key is the destination and the value
24         is a RoutingTableEntry. e.g. {destination: RoutingTableEntry}
25         A RoutingTableEntry contains info about the next_hop, metric, and timeouts.
26     """
27
28     def __init__(self, debug_func, config, output_socket):
29         global debug
30         debug = debug_func
31         self.our_routerid = config.router_id
32         self.output_routers = config.outputs
33         self.socket = output_socket
34
35         self.routing_table = {}
36         self.next_periodic_update = time.time()
37         self.triggered_update_pending = False
38         self.next_triggered_update = 0
39
40
41     def __str__(self):
42         lines = f'''Router {self.our_routerid:<16} Routing Table
43 +-----+-----+-----+-----+
44 | destination | next hop | metric | update due | deletion due |
45 +-----+-----+-----+-----+
46 '''
47
48         for dest, entry in sorted(self.routing_table.items()):
49             deletion_due = entry.deletion_due_in()
50             if deletion_due == math.inf:
51                 deletion_due = ''
52             else:
```

```
52             deletion_due = int(deletion_due)
53             lines += f'| {dest:>11} | {entry.next_hop:>8} | {entry.metric:>6} '
54             lines += f'| {entry.update_due_in():>10.0f} | {deletion_due:>12} |\n'
55             lines += '+-----+-----+-----+-----+-----+\n'
56         return lines
57
58
59     def table_list(self):
60         """Return a list of routing table entries. Does not include
61         timeout times, but does include a deletion process flag.
62         Used for automatic testing and to detect routing table changes.
63         """
64         return [[d, e.next_hop, e.metric, e.deletion_process_underway()] for d,e in sorted(self.routing_table.items())]
65
66
67     def next_timeout(self):
68         """Return the time in seconds (as a float) until the next timeout.
69         Timeouts include periodic update messages (every 30 seconds), and
70         triggered updates related to routing table entries. Triggered
71         updates occur if a routing table entry hasn't been updated for
72         180 seconds, or if a routing table entry has been garbage
73         collected for 120 seconds.
74         """
75         next_periodic_update_in = self.next_periodic_update - time.time()
76         next_periodic_update_in = max(0, next_periodic_update_in)
77
78         timeouts = [next_periodic_update_in]
79         for entry in self.routing_table.values():
80             timeouts.append(entry.next_timeout())
81
82         if self.triggered_update_pending:
83             next_triggered_update_in = self.next_triggered_update - time.time()
84             timeouts.append(next_triggered_update_in)
85
86         smallest_timeout = min(timeouts)
87         return max(0, smallest_timeout)
88
89
90     def incoming_message(self, message):
91         """Process an incoming UDP packet."""
92         try:
93             rip_packet = RipPacket(message)
94         except AssertionError as e:
95             debug(f'Received invalid packet: {e}')
96             return
97
98         next_hop = rip_packet.routerid
99         if next_hop not in self.output_routers:
100             debug(f'Received packet from unknown router {next_hop}')
101             return
```

```
102     _, metric_to_next_hop = self.output_routers[next_hop]
103     self.add_to_table(next_hop, next_hop, metric_to_next_hop) # add sender to ↵
104     routing table
105
106     for rip_entry in rip_packet.entries:
107         metric = min(metric_to_next_hop + rip_entry.metric, INFINITE_METRIC)
108         self.add_to_table(rip_entry.routerid, next_hop, metric)
109
110 def add_to_table(self, destination, next_hop, metric):
111     """Update or add a table entry.
112     Only add a new entry if the metric isn't infinity.
113     The RIP assignment says to not send a triggered message for
114     metric updates or new routes.
115     """
116
117     if destination == self.our_routerid:
118         return # don't add ourself to our routing table
119     if destination in self.routing_table.keys():
120         reason = self.routing_table[destination].update_entry(next_hop, metric)
121         if reason:
122             debug(f'{self.our_routerid} updating routing table entry for ↵
123             destination {destination}:')
124             debug(f'    {reason}')
125         elif metric < INFINITE_METRIC:
126             debug(f'{self.our_routerid} added a new route to destination {destination} ↵
127             next-hop {next_hop} metric {metric}')
128             self.routing_table[destination] = RoutingTableEntry(next_hop, metric)
129
130
131 def send_any_updates(self):
132     """Check if a periodic or triggered update should be sent.
133     Triggered updates only for when routes become invalid (route
134     deleted or metric set to 16), not for new/updated routes.
135     After sending a triggered update, don't send future triggered
136     updates for 1 to 5 seconds.
137     """
138
139     to_delete = []
140     for destination, entry in self.routing_table.items():
141         if entry.should_delete():
142             to_delete.append(destination)
143             self.triggered_update_pending = True
144         elif entry.should_begin_deletion():
145             debug(f'Starting deletion process for destination {destination}')
146             entry.begin_deletion()
147             self.triggered_update_pending = True
148
149     for dest in to_delete: # since you can't delete entries while iterating ↵
150         over them
151         debug(f'Deleting destination {dest}')
152         del self.routing_table[dest]
```

```
149         periodic_update = time.time() >= self.next_periodic_update
150         triggered_update = self.triggered_update_pending and time.time() >= self.
151             next_triggered_update
152         if periodic_update or triggered_update:
153             self.send_response_messages()
154
155     def send_response_messages(self):
156         """Send a periodic/triggered update message.
157         Send a response message to all neighbours
158         containing the complete routing table (as set by assignment
159         specifications) utilising split-horizon with poisoned-reverse.
160         The next periodic update message should be sent in
161         30 seconds +/- up to 5 seconds (1/6th of 30s) randomly.
162         The next triggered update message should be sent in
163         1 (1/5th of 5 seconds) to 5 seconds randomly.
164         """
165
166         for router_id, [port, metric] in self.output_routers.items():
167             packets = self.build_packets(router_id)
168             for p in packets:
169                 try:
170                     RipPacket(p)
171                 except AssertionError as e:
172                     debug(f'Sending invalid packet: {e}')
173                     self.socket.sendto(p, ('127.0.0.1', port))
174
175             self.next_periodic_update = (time.time() +
176                 PERIODIC_UPDATE_DELAY +
177                 random.uniform(-PERIODIC_UPDATE_DELAY/6, PERIODIC_UPDATE_DELAY/6))
178             self.triggered_update_pending = False
179             self.next_triggered_update = (time.time() +
180                 random.uniform(TRIGGERED_UPDATE_DELAY/5, TRIGGERED_UPDATE_DELAY))
181
182     def build_packets(self, destination_router_id):
183         """Return response message packets to be sent to the defined
184         router. Utilises split-horizon with optional poisoned-reverse.
185         """
186         packets = []
187
188         packet = self.empty_rip_packet()
189         packet += rip_entry(destination_router_id, INFINITE_METRIC) # always add
190             the receiver as a rip entry with inf metric
191
192         for destination, entry in self.routing_table.items():
193             metric = entry.metric
194             if entry.next_hop == destination_router_id:
195                 if POISONED_REVERSE:
196                     metric = INFINITE_METRIC
197                 else:
198                     continue # don't add the entry
```

```
198
199         if len(packet) >= (4 + 20*25): # if 25 entries
200             packets.append(packet)
201             packet = self.empty_rip_packet()
202
203             packet += rip_entry(destination, metric)
204
205             packets.append(packet)
206             return packets
207
208
209     def empty_rip_packet(self):
210         """Return an empty rip packet (headers only).
211         RFC all-zeros field is used for the routerid by assignment specs.
212         """
213
214         packet = bytearray(4)
215         packet[0] = 2 # command
216         packet[1] = 2 # version
217         packet[2:4] = self.our_routerid.to_bytes(2)
218         return packet
219
220
221     def rip_entry(destination, metric):
222         """Return a rip entry for use in a rip packet."""
223         entry = bytearray(20)
224         entry[0:2] = (2).to_bytes(2) # address family identifier
225         entry[4:8] = destination.to_bytes(4)
226         entry[16:20] = metric.to_bytes(4)
227         return entry
228
229
230     class RoutingTableEntry:
231         """A single entry for use in the routing table.
232         The RFC's 'garbage-collection' is called 'deletion' here.
233         Route change flags are not used due to us not sending triggered
234         updates for route metric changes according to the RIP assignment.
235         """
236
237         def __init__(self, next_hop, metric):
238             self.next_hop = next_hop
239             self.metric = metric
240             self.time_update_due = time.time() + ENTRY_TIMEOUT_DELAY
241             self.time_deletion_due = None
242
243         def deletion_process_underway(self):
244             return self.time_deletion_due != None
245
246
247         def over_halfway_to_update_due(self):
248             due_in = self.time_update_due - time.time()
```

```
249         return due_in <= ENTRY_TIMEOUT_DELAY/2
250
251
252     def update_due_in(self):
253         """Time in seconds until an update is due."""
254         due_in = self.time_update_due - time.time()
255         return max(0, due_in)
256
257
258     def deletion_due_in(self):
259         """Time in seconds until deletion is due."""
260         due_in = math.inf
261         if self.deletion_process_underway():
262             due_in = self.time_deletion_due - time.time()
263         return max(0, due_in)
264
265
266     def next_timeout(self):
267         """Return the time in seconds (as a float) until the next timeout."""
268         smallest_time = min(self.update_due_in(), self.deletion_due_in())
269         return max(0, smallest_time)
270
271
272     def update_entry(self, next_hop, new_metric):
273         """If the deletion process is underway for a route, replace it.
274         If the new metric is 16 then don't add it (no better than current).
275         Return a string describing the reason for change.
276         """
277
278         reason = None
279         update_timeouts = False
280
281         if next_hop == self.next_hop:
282             update_timeouts = True
283             if self.metric != new_metric:
284                 reason = f'updated next-hop {self.next_hop} metric from {self.metric} to {new_metric} (update is from next-hop)'
285                 self.metric = new_metric
286
287             elif new_metric < self.metric:
288                 reason = f'updated next-hop from {self.next_hop} ({self.metric}) to {next_hop} ({new_metric}) (better metric)'
289                 update_timeouts = True
290                 self.next_hop = next_hop
291                 self.metric = new_metric
292
293             # RFC section 3.9.2 heuristic
294             elif (new_metric != INFINITE_METRIC and
295                  new_metric == self.metric and
296                  self.over_halfway_to_update_due()):
297                 update_timeouts = True
298                 reason = f'updated next-hop from {self.next_hop} ({self.metric}) to {next_hop} ({new_metric}) (over halfway to update due)'
```

```
298         next_hop} ({new_metric}) (over halfway to update due)'  
299         self.next_hop = next_hop  
300         self.metric = new_metric  
301  
302     if update_timeouts:  
303         self.time_update_due = time.time() + ENTRY_TIMEOUT_DELAY  
304     if self.metric < INFINITE_METRIC:  
305         self.time_deletion_due = None  
306  
307     return reason  
308  
309 def should_begin_deletion(self):  
310     """Return True if the deletion process should be started.  
311     Deletion process should not be started if it is already underway.  
312     """  
313     if not self.deletion_process_underway():  
314         return (self.metric >= INFINITE_METRIC or  
315                 time.time() >= self.time_update_due)  
316     return False  
317  
318  
319 def begin_deletion(self):  
320     assert self.deletion_process_underway() is False  
321     self.metric = INFINITE_METRIC  
322     self.time_deletion_due = time.time() + GARBAGE_COLLECTION_DELAY  
323  
324  
325 def should_delete(self):  
326     """Return True if this entry should be deleted immediately."""  
327     if self.deletion_process_underway():  
328         return time.time() >= self.time_deletion_due  
329     return False  
330  
331  
332 class RipPacket:  
333     """This class represents a validated RIP request packet.  
334     If a RIP packet entry is invalid, ignore it.  
335     1 byte - command (must be 2)  
336     1 byte - version (must be 2)  
337     2 bytes - routerid (all-zeros in RIP RFC)  
338     20 bytes - rip entry (1 to 25 lots of these)  
339     """  
340  
341     def __init__(self, packet):  
342         self.validate_rip_packet(packet)  
343         self.routerid = int.from_bytes(packet[2:4])  
344         self.entries = []  
345         for i in range(4, len(packet), 20):  
346             try:  
347                 self.entries.append(RipEntry(packet[i: i+20]))  
348             except AssertionError as e:
```

```
348                 debug(f'RIP packet entry error: {e}')
349
350     def __str__(self):
351         lines = f'''packet:
352 Source: {self.routerid}'''
353         for entry in self.entries:
354             lines += f"""
355 {entry}"""
356         if not self.entries:
357             lines += f"""
358 <EMPTY PACKET>"""
359
360         return lines + '\n'
361
362     def validate_rip_packet(self, packet):
363         """Raise an AssertionError if the packet is invalid.
364         Does not check the validity of the contained rip entries.
365         """
366         assert len(packet) >= 4+20, f"packet length invalid: {len(packet)}"
367         assert len(packet) <= 4+20*25, f"packet length invalid: {len(packet)}"
368         assert (len(packet) - 4) % 20 == 0, f"packet length invalid: {len(packet)}"
369         assert packet[0] == 2, "command field not 2"
370         assert packet[1] == 2, "version field not 2"
371         routerid = int.from_bytes(packet[2:4])
372         assert routerid_is_valid(routerid), f"router-id invalid {routerid}"
373
374
375 class RipEntry:
376     """This class represents a validated RIP entry from a RIP packet.
377     2 bytes - address family (ignore)
378     2 bytes - all zeros
379     4 bytes - routerid (IPv4 in RIP RFC)
380     8 bytes - all zeros
381     4 bytes - metric
382     """
383
384     def __init__(self, entry):
385         self.validate_rip_entry(entry)
386         self.routerid = int.from_bytes(entry[4:8])
387         self.metric = int.from_bytes(entry[16:20])
388
389     def __str__(self):
390         return f'router-id: {self.routerid} metric: {self.metric}'
391
392     def validate_rip_entry(self, entry):
393         """Raise an AssertionError if the rip entry is invalid."""
394         assert len(entry) == 20, "RIP entry length not 20"
395         assert int.from_bytes(entry[0:2]) == 2, "address family must be 2"
396         assert int.from_bytes(entry[2:4]) == 0, "field must be all zeros"
397         routerid = int.from_bytes(entry[4:8])
398         assert routerid_is_valid(routerid), f"router-id invalid {routerid}"
399         assert int.from_bytes(entry[8:16]) == 0, "field must be all zeros"
400         metric = int.from_bytes(entry[16:20])
```

```
399     assert metric_is_valid(metric), f"metric invalid {metric}"  
400
```

```
1 import configparser
2 from itertools import combinations
3
4
5 class Config:
6     def __init__(self, router_id, input_ports, outputs):
7         self.router_id = router_id
8         self.input_ports = input_ports
9         self.outputs = outputs
10
11    def __str__(self):
12        lines = f"""CONFIG:
13        router id: {self.router_id}
14        input ports: {self.input_ports}
15        outputs:"""
16        for routerid, [port, metric] in self.outputs.items():
17            lines += f"""
18            router-id: {routerid} port: {port} metric: {metric}"""
19        return lines + '\n'
20
21
22    def read_config_file(filename):
23        config = configparser.ConfigParser()
24        config.read(filename)
25        try:
26            return get_config(config)
27        except ValueError as e:
28            raise ValueError(f'CONFIG {filename} ERROR: {e}')
29
30
31    def get_config(config):
32        """
33        >>> config = configparser.ConfigParser()
34        >>> config['SETTINGS'] =
35        {'router-id': '2', 'input-ports': '2000', 'outputs': '3000-1-3'}
36        >>> c1 = get_config(config)
37        >>> print(c1)
38        CONFIG:
39            router id: 2
40            input ports: [2000]
41            outputs:
42                router-id: 3 port: 3000 metric: 1
43            <BLANKLINE>
44
45            """
46        router_id, input_ports, outputs = validate_config(config)
47        return Config(router_id, input_ports, outputs)
48
49    def validate_configs_by_filename(filenames):
50        configs = [read_config_file(filename) for filename in filenames]
```

```
51     validate_configs(configs)
52
53 def validate_configs(configs):
54     """For all the provided configs:
55     ensures that all router-ids are unique,
56     sending/receiving router-ids match between neighbours,
57     and that metrics between neighbours are the same.
58
59     >>> config1 = configparser.ConfigParser()
60     >>> config1['SETTINGS'] =
61         {'router-id':'2','input-ports':'2000','outputs':'3000-1-3'}
62     >>> config2 = configparser.ConfigParser()
63     >>> config3 = configparser.ConfigParser()
64
65     >>> config2['SETTINGS'] =
66         {'router-id':'3','input-ports':'3000','outputs':'2000-1-2'}
67     >>> validate_configs([get_config(config1), get_config(config2)])
68
69     >>> config2['SETTINGS'] =
70         {'router-id':'2','input-ports':'3000','outputs':'2000-1-2'}
71     >>> validate_configs([get_config(config1), get_config(config2)])
72     Traceback (most recent call last):
73     AssertionError: same router-id: 2
74
75     >>> config2['SETTINGS'] =
76         {'router-id':'3','input-ports':'3333','outputs':'3000-1-2'}
77     >>> validate_configs([get_config(config1), get_config(config2)])
78     Traceback (most recent call last):
79     AssertionError: port 3000 is already an output to router 3
80
81     >>> config2['SETTINGS'] =
82         {'router-id':'3','input-ports':'3000','outputs':'2000-1-3'}
83     >>> validate_configs([get_config(config1), get_config(config2)])
84     Traceback (most recent call last):
85     AssertionError: router-id mismatch between routers 2 and 3 on port 2000
86
87     >>> config2['SETTINGS'] =
88         {'router-id':'3','input-ports':'3000','outputs':'2222-1-2'}
89     >>> validate_configs([get_config(config1), get_config(config2)])
90     Traceback (most recent call last):
91     AssertionError: router 2 listening on port 2000 but no sender
92
93     >>> config2['SETTINGS'] =
94         {'router-id':'3','input-ports':'3333','outputs':'2000-1-2'}
95     >>> validate_configs([get_config(config1), get_config(config2)])
96     Traceback (most recent call last):
97     AssertionError: sending to router 3 on port 3000 but no receiver
98
99     >>> config2['SETTINGS'] =
100        {'router-id':'3','input-ports':'3000','outputs':'2000-2-2'}
101    >>> validate_configs([get_config(config1), get_config(config2)])
```

```

94     Traceback (most recent call last):
95     AssertionError: metric mismatch between routers 2 and 3
96
97     >>> config1['SETTINGS'] =
98     {'router-id':'2','input-ports':'2000,2001','outputs':'3000-1-3,4000-2-4'}
99     >>> config2['SETTINGS'] =
100     {'router-id':'3','input-ports':'3000,3001','outputs':'2000-1-2,4001-3-4'}
101     >>> config3['SETTINGS'] =
102     {'router-id':'4','input-ports':'4000,4001','outputs':'2001-2-2,3001-3-3'}
103     >>> validate_configs([get_config(config1), get_config(config2),
104                           get_config(config3)])
105     """
106     for c1, c2 in combinations(configs, 2):
107         assert c1.router_id != c2.router_id, f'same router-id: {c1.router_id}'
108
109     port_ids = {} # {port: [input_id, output_id]}
110     metrics = {} # {(router1_id, router2_id), metric} # where router1_id <
111     router2_id
112     for config in configs:
113         for port in config.input_ports:
114             current_ids = port_ids.get(port, [None, None])
115             assert current_ids[0] is None, f'port {port} already an input for
116             router {current_ids[0]}'
117             current_ids[0] = config.router_id
118             port_ids[port] = current_ids
119
120             for router_id, [port, metric] in config.outputs.items():
121                 current_ids = port_ids.get(port, [None, None])
122                 assert current_ids[1] is None, f'port {port} is already an output to
123                 router {current_ids[1]}'
124                 current_ids[1] = router_id
125                 port_ids[port] = current_ids
126
127                 lower_id, upper_id = sorted([config.router_id, router_id])
128                 current_metric = metrics.get((lower_id, upper_id), None)
129                 if current_metric is not None:
130                     assert current_metric == metric, f'metric mismatch between routers
131                     {lower_id} and {upper_id}'
132                     metrics[(lower_id, upper_id)] = metric
133
134             for port, [in_id, out_id] in port_ids.items():
135                 assert in_id != None, f'sending to router {out_id} on port {port} but no
136                 receiver'
137                 assert out_id != None, f'router {in_id} listening on port {port} but no
138                 sender'
139                 assert in_id == out_id, f'router-id mismatch between routers {in_id} and {
140                 out_id} on port {port}'
141
142     def routerid_is_valid(routerid):
143         return 1 <= routerid <= 64000

```

```
134
135     def validate_router_id(routerid):
136         """
137             >>> validate_router_id('1')
138                 1
139             >>> validate_router_id('64000')
140                 64000
141             >>> validate_router_id('0')
142                 Traceback (most recent call last):
143                 ValueError: router-id must be a number between 1 and 64000. Got "0"
144             >>> validate_router_id('64001')
145                 Traceback (most recent call last):
146                 ValueError: router-id must be a number between 1 and 64000. Got "64001"
147                 """
148             routerid = routerid.strip()
149             if routerid.isdigit() and routerid_is_valid(int(routerid)):
150                 return int(routerid)
151             else:
152                 raise ValueError(f'router-id must be a number between 1 and 64000. Got "{routerid}"')
153
154
155     def port_is_valid(port):
156         return 1024 <= port <= 64000
157
158     def validate_port(port):
159         """
160             >>> validate_port('1024')
161                 1024
162             >>> validate_port('64000')
163                 64000
164             >>> validate_port('1023')
165                 Traceback (most recent call last):
166                 ValueError: port must be a number between 1024 and 64000. Got "1023"
167             >>> validate_port('64001')
168                 Traceback (most recent call last):
169                 ValueError: port must be a number between 1024 and 64000. Got "64001"
170                 """
171             port = port.strip()
172             if port.isdigit() and port_is_valid(int(port)):
173                 return int(port)
174             else:
175                 raise ValueError(f'port must be a number between 1024 and 64000. Got "{port}"')
176
177
178     def metric_is_valid(metric):
179         return 1 <= metric <= 16
180
181     def validate_metric(metric):
182         """
```

```
183     >>> validate_metric('1')
184     1
185     >>> validate_metric('16')
186     16
187     >>> validate_metric('0')
188     Traceback (most recent call last):
189     ValueError: metric must be a number between 1 and 16. Got "0"
190     >>> validate_metric('17')
191     Traceback (most recent call last):
192     ValueError: metric must be a number between 1 and 16. Got "17"
193     """
194     metric = metric.strip()
195     if metric.isdigit() and metric_is_valid(int(metric)):
196         return int(metric)
197     else:
198         raise ValueError(f'metric must be a number between 1 and 16. Got "{metric}"')
199
200
201 def validate_config(config):
202     """
203     >>> config = configparser.ConfigParser()
204     >>> validate_config(config)
205     Traceback (most recent call last):
206     ValueError: SETTINGS header not found
207
208     >>> config['SETTINGS'] = {'input-ports':'1024','outputs':'64000-0-1'}
209     >>> validate_config(config)
210     Traceback (most recent call last):
211     ValueError: "router-id" parameter not found
212
213     >>> config['SETTINGS'] = {'router-id':'1','outputs':'64000-0-1'}
214     >>> validate_config(config)
215     Traceback (most recent call last):
216     ValueError: "input-ports" parameter not found
217
218     >>> config['SETTINGS'] = {'router-id':'1','input-ports':'1024'}
219     >>> validate_config(config)
220     Traceback (most recent call last):
221     ValueError: "outputs" parameter not found
222
223     >>> config['SETTINGS'] =
224     {'router-id':'1','input-ports':'2000,2000','outputs':'5000-15-1'}
225     >>> validate_config(config)
226     Traceback (most recent call last):
227     ValueError: "2000" is a duplicate port number
228
229     >>> config['SETTINGS'] =
230     {'router-id':'1','input-ports':'2000','outputs':'2000-15-1'}
231     >>> validate_config(config)
232     Traceback (most recent call last):
```

```
231     ValueError: "2000" is already defined as an input port
232
233     >>> config['SETTINGS'] =
234         {'router-id':'1','input-ports':'1024','outputs':'64000-1-1'}
235     >>> validate_config(config)
236     (1, [1024], {1: [64000, 1]})

237     >>> config['SETTINGS'] = {'router-id':' 01 ','input-ports':' 01024 ,
238         01025','outputs':' 064000 - 011 - 01 , 05000 - 012 - 02'}
239     >>> validate_config(config)
240     (1, [1024, 1025], {1: [64000, 11], 2: [5000, 12]})

241     >>> config['SETTINGS'] =
242         {'router-id':'1','input-ports':'2000,2001,2002','outputs':'5000-14-2,5001-15-64
243         000'}
244     >>> validate_config(config)
245     (1, [2000, 2001, 2002], {2: [5000, 14], 64000: [5001, 15]})

246     """
247     if not 'SETTINGS' in config:
248         raise ValueError('SETTINGS header not found')
249     for param in ['router-id', 'input-ports', 'outputs']:
250         if not param in config['SETTINGS']:
251             raise ValueError(f'{param}' parameter not found')

252     router_id = config['SETTINGS']['router-id']
253     router_id = validate_router_id(router_id)

254
255     input_ports_str = config['SETTINGS']['input-ports'].split(',')
256     input_ports = []
257     for port in input_ports_str:
258         port = validate_port(port)
259         if port in input_ports:
260             raise ValueError(f'{port}' is a duplicate port number)
261         else:
262             input_ports.append(port)

263
264     outputs_str = config['SETTINGS']['outputs'].split(',')
265     outputs = {}
266     for output in outputs_str:
267         port, metric, out_routerid = output.strip().split('-')

268         port = validate_port(port)
269         if port in input_ports:
270             raise ValueError(f'{port}' is already defined as an input port)
271         metric = validate_metric(metric)
272         out_routerid = validate_router_id(out_routerid)

273         outputs[out_routerid] = [port, metric]

274
275     if input_ports == []:
```

```
278         raise ValueError(f'There must be at least one input port')
279     if outputs == []:
280         raise ValueError(f'There must be at least one output')
281
282     return router_id, input_ports, outputs
283
284
285 if __name__ == '__main__':
286     import doctest
287     results = doctest.testmod()
288     print(results)
289
```

```
1 import configparser
2 import fcntl
3 from itertools import combinations
4 import json
5 import math
6 import os
7 import random
8 import selectors
9 from subprocess import Popen, PIPE, STDOUT
10 import time
11
12 from configmanager import validate_configs_by_filename
13
14
15 NUM_ROUTERS = 100
16
17
18 FOLDER = 'test_configs'
19 os.makedirs(FOLDER, exist_ok=True)
20
21
22 class Test:
23     def __init__(self, neighbour_func, change_topology=None, topology_changes=1):
24         self.make_neighbours_func = neighbour_func
25         self.topology_change_func = change_topology
26         self.topology_changes_remaining = topology_changes
27
28     def make_neighbours(self, processes):
29         self.make_neighbours_func(processes)
30
31     def can_change_topology(self):
32         return self.topology_change_func != None and self.
33             topology_changes_remaining > 0
34
35     def change_topology(self, processes):
36         self.topology_changes_remaining -= 1
37         self.topology_change_func(processes)
38
39 ports = iter(range(10000, 64000))
40 def make_neighbours(p1, p2):
41     port1 = next(ports)
42     port2 = next(ports)
43     metric = random.randint(1, 15)
44     p1.add_neighbour(port1, port2, metric, p2)
45     p2.add_neighbour(port2, port1, metric, p1)
46
47 def fully_connected(processes):
48     for p1, p2 in combinations(processes, 2):
49         make_neighbours(p1, p2)
50
```

```
51 def sparsely_connected(processes):
52     rand_processes = list(processes)
53     random.shuffle(rand_processes)
54     for p1 in processes:
55         num_neighbours = 0
56         for p2 in rand_processes:
57             if p1.routerid != p2.routerid and p2.routerid not in p1.get_neighbours():
58                 make_neighbours(p1, p2)
59                 num_neighbours += 1
60                 if num_neighbours >= 1:
61                     break
62
63 def change_topology(processes):
64     processes = list(processes)
65     if random.choice([False, True]):
66         to_stop = random.sample(processes, len(processes)//2)
67         print(f'stopping {len(to_stop)} processes randomly')
68         for p in to_stop:
69             p.stop()
70     else:
71         to_start = random.sample(processes, len(processes)//2)
72         print(f'starting {len(to_start)} processes randomly')
73         for p in to_start:
74             p.start()
75
76 test1 = Test(fully_connected)
77
78 test2 = Test(sparsely_connected)
79
80 test3 = Test(fully_connected, change_topology, 5)
81
82 test4 = Test(sparsely_connected, change_topology, 10)
83
84
85 class ProcessManager:
86     def __init__(self):
87         self.processes_dict = {}
88
89     def get_processes(self):
90         return self.processes_dict.values()
91
92     def get_alive_processes(self):
93         return [p for p in self.processes_dict.values() if p.alive]
94
95     def get_process(self, id):
96         return self.processes_dict[id]
97
98     def start_processes(self):
99         for p in self.get_processes():
100            p.start()
```

```
101
102     def stop_processes(self):
103         for p in self.get_processes():
104             p.stop()
105
106     def new_processes(self):
107         self.stop_processes()
108         for i in range(1, NUM_ROUTERS+1):
109             self.processes_dict[i] = Process(i)
110
111     def setup_test(self, test):
112         self.new_processes()
113         test.make_neighbours(self.get_processes())
114         self.write_configs()
115         validate_configs_by_filename([p.filename for p in self.get_processes()])
116         self.start_processes()
117
118     def change_test_topology(self, test):
119         test.change_topology(self.get_processes())
120         for p in self.get_processes():
121             p.clear_routing_table()
122
123     def write_configs(self):
124         for p in self.get_processes():
125             p.write_config()
126
127
128 class Process:
129     def __init__(self, routerid):
130         self.routerid = routerid
131         self.inputs = []
132         self.outputs = {}
133         self.filename = f'{FOLDER}/autoconfig{self.routerid}.ini'
134         self.process = None
135         self.alive = False
136
137         self.routing_table = None
138         self.routing_table_time = math.inf
139         self.have_checked_convergence = False
140         self.converged = False
141
142
143     def __str__(self):
144         return str(self.routerid)
145
146
147     def add_neighbour(self, in_port, out_port, metric, neighbour):
148         self.inputs.append(str(in_port))
149         self.outputs[neighbour.routerid] = [neighbour, out_port, metric]
150
151
```

```
152     def get_neighbours(self):
153         return self.outputs
154
155
156     def write_config(self):
157         config = configparser.ConfigParser()
158         config['SETTINGS'] = {
159             'router-id': str(self.routerid),
160             'input-ports': ','.join(self.inputs),
161             'outputs': ','.join(f'{port}-{metric}-{id}' for id, _, port, metric) ↵
162             in self.outputs.items())
163         }
164         with open(self.filename, 'w') as file:
165             config.write(file)
166
167
168     def start(self):
169         """Start the process and make its stdout non-blocking."""
170         if not self.alive:
171             self.alive = True
172             self.process = Popen(["python", "daemon.py", self.filename, ↵
173             "--autotesting"], stdout=PIPE, stderr=STDOUT)
174             fcntl.fcntl(self.process.stdout.fileno(), fcntl.F_SETFL, os.O_NONBLOCK)
175
176     def stop(self):
177         self.alive = False
178         self.process.kill()
179
180
181     def get_stdout(self):
182         return self.process.stdout
183
184
185     def read_line(self):
186         line = self.process.stdout.readline()
187         if line:
188             line = line.decode().strip()
189             try:
190                 line = json.loads(line)
191             except json.decoder.JSONDecodeError as e:
192                 print(self, 'decode error', line)
193                 return
194             if type(line) != list:
195                 print(self, 'received non-list', line)
196                 return
197
198             if line != self.routing_table:
199                 self.routing_table = line
200                 self.routing_table_time = time.time()
201                 self.have_checked_convergence = False
```

```
201             self.converged = False
202
203
204     def clear_routing_table(self):
205         self.routing_table = None
206         self.routing_table_time = math.inf
207         self.have_checked_convergence = False
208         self.converged = False
209
210
211     def routing_table_entries(self):
212         return {routerid:metric for routerid, _, metric, _ in self.routing_table}
213
214
215     def check_convergence(self):
216         # an offline router is considered converged
217         if not self.alive:
218             self.converged = True
219             return
220
221         # don't check for convergence again if the routing table hasn't changed
222         if self.have_checked_convergence:
223             return
224
225         # only check if routing table hasn't changed for 10 seconds
226         if time.time() - self.routing_table_time < 10:
227             return
228
229         self.calculate_convergence()
230
231
232     def calculate_convergence(self):
233         min_costs, parents = dijkstras(self.routerid)
234         routing_table_entries = self.routing_table_entries()
235
236         self.converged = True
237         for routerid, metric in min_costs.items():
238             if metric >= 16 or routerid == self.routerid:
239                 continue
240
241             if routerid not in routing_table_entries:
242                 self.converged = False
243                 print(f'{self} not converged to router {routerid} (not in routing table, cost should be: {metric})')
244                 print('Dijkstras path:', dijkstras_path(min_costs, parents, self.routerid, routerid))
245                 print()
246                 continue
247
248             actual_metric = routing_table_entries[routerid]
249             if actual_metric != metric:
```

```

250                     self.converged = False
251                     print(f'{self} not converged to router {routerid} (current cost: { ↵
252                         actual_metric}, should be: {metric})')
253                     print('Dijkstras path:', dijsktras_path(min_costs, parents, self. ↵
254                         routerid, routerid))
255                     print('Current path: ', end='')
256                     print_actual_path(self.routerid, routerid)
257                     print()
258
259                     self.have_checked_convergence = True
260
261
262
263
264     def dijkstras(source_id):
265         dist = {}
266         prev = {}
267         queue = []
268         for p in processmanager.get_alive_processes():
269             id = p.routerid
270             dist[id] = math.inf
271             prev[id] = None
272             queue.append(id)
273         assert source_id in dist
274         dist[source_id] = 0
275
276         while queue:
277             u = None
278             min_dist = math.inf
279             for v in queue:
280                 if dist[v] <= min_dist:
281                     u = v
282                     min_dist = dist[v]
283             queue.remove(u)
284
285             u_neighbours = processmanager.get_process(u).get_neighbours()
286             for v, [process, _, metric] in u_neighbours.items():
287                 if v not in queue:
288                     continue
289
290                     cost = dist[u] + metric
291                     if cost <= dist[v]:
292                         dist[v] = cost
293                         prev[v] = u
294
295         return dist, prev
296
297
298     def dijsktras_path(dist, prev, src, dest):
299         current = dest
300         path = f'{current} ({dist[current]})'
301         while current != src:
302             current = prev[current]

```

```
299         path = f'{current} ({dist[current]}) --> ' + path
300     return path
301
302
303 def print_actual_path(src, dest, depth=0):
304     if depth > 15:
305         print('ABORTING')
306         return
307     if src == dest:
308         print(f'{src} (0)')
309         return
310
311     src_routing_table = processmanager.get_process(src).routing_table
312     if src_routing_table == None:
313         print(f'{src} (no route to {dest})')
314         return
315     for routerid, nexthop, metric, _ in src_routing_table:
316         if routerid == dest:
317             break
318     print(f'{src} ({metric}) --> ', end=' ')
319     print_actual_path(nexthop, dest, depth+1)
320
321
322 processmanager = ProcessManager()
323
324 def main():
325     tests = [test1, test2, test3, test4]
326     for i in range(len(tests)):
327         test = tests[i]
328         processmanager.setup_test(test)
329         print(f'test {i} starting')
330         run_to_convergence()
331         while test.can_change_topology():
332             print(f'test {i} changing topology')
333             processmanager.change_test_topology(test)
334             run_to_convergence()
335         print(f'test {i} finished')
336
337
338 def run_to_convergence():
339     selector = selectors.DefaultSelector()
340     for p in processmanager.get_processes():
341         selector.register(p.get_stdout(), selectors.EVENT_READ, p)
342
343     prev_not_converged = []
344     while True:
345         events = selector.select(timeout=1)
346         for key, _ in events:
347             p = key.data
348             p.read_line()
```

```
350     all_converged = True
351     not_converged = []
352     for p in processmanager.get_processes():
353         p.check_convergence()
354         if not p.converged:
355             all_converged = False
356             not_converged.append(p.routerid)
357
358     if all_converged:
359         print('all routers converged correctly')
360         return
361     elif not_converged != prev_not_converged:
362         prev_not_converged = not_converged
363         print(len(not_converged), 'routers not converged.', not_converged[:10])
364
365
366 try:
367     main()
368 except KeyboardInterrupt:
369     pass
370 finally:
371     processmanager.stop_processes()
372     print('exiting')
373
```